Christopher Jones

DRAFT
Achieving distributed consensus with Paxos

Part II Project

Trinity Hall

May 13, 2018

Proforma

Name:	**Christopher Jones**
College:	**Trinity Hall**
Project Title:	**Achieving distributed consensus with Paxos**
Examination:	**Computer Science Tripos — Part II, June 2018**
Word Count:	...[1]
Project Originator:	Christopher Jones
Supervisor:	Dr Richard Mortier

Original Aims of the Project

...

Work Completed

...

Special Difficulties

None.

---

[1]This word was computed by `detex diss.tex | tr -cd '0-9A-Za-z \n' | wc -w`

Declaration

I, [Name] of [College], being a candidate for Part II of the Computer Science Tripos [or the Diploma in Computer Science], hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

Signed [signature]

Date [date]

Contents

List of Figures

Chapter 1

Introduction

This dissertation describes the process of researching, developing and evaluating an implementation of the Multi–Paxos distributed consensus algorithm in OCaml.

1.1   Background

Distributed systems suffer from a number of possible errors and failure modes. Unreliability is present in the network where messages can be delayed, re–ordered and dropped and processes can exhibit faulty behaviour such as stalling and crashing. The result of this is that distributed systems can end up inconsistent states and even unable to make progress.

Consensus is the reaching of agreement in the face of such unreliable conditions. Applications such as transaction systems and distributed databases require consensus in order to remain consistent. Consensus algorithms provide a means of reaching agreement in the face of such unreliability; this is crucial in the design of distributed systems.

Paxos is a consensus algorithm first described by Lamport [5] that allows for consensus to be reached under the typical unreliable conditions present in a distributed system. The algorithm relies on processes participating in a voting protocol that tolerates the failure of a minority of processes.

Paxos is used internally in large–scale production systems such as Google's Chubby [1] distributed lock service, where it is used to maintain consistency between replicas. Microsoft's Autopilot [4] system for data centre management also uses Paxos, again to replicate data across machines. The generality of Paxos allows it to be used as an underlying primitive for various distributed systems techniques. State Machine Replication [12] is a technique whereby any application that behaves like a state machine can be replicated across a number of machines participating in the Paxos protocol. Likewise,

atomic broadcast [11] can be implemented with Paxos as an underlying primitive.

Over time Paxos has been extended and modified to emphasise different performance trade–offs. Multi–Paxos is the most typically deployed variant which allows for explicit agreement over a sequence of values. Another example, Fast Paxos [7], is a variant that reduces the number of message delays between proposing a value and receiving a response. More recently, Flexible Paxos [3] is a variant that relaxes the requirement on agreement between participants in the synod protocol.

There are also a number of similar algorithms that are not based on Paxos. View-stamped replication [9] is primarily a protocol for implementing state machine replication that was developed independently of Paxos. Raft [10] is a modern alternative to Paxos that attempts to be more understandable.

1.2   Aims

The aim of this project was to produce an implementation of the Multi–Paxos variant of the Paxos algorithm. This is the variant that is used most widely in production systems and provides a foundation upon which to perform state machine replication; in this case to replicate a key value store application.

From the outset OCaml was to be used as the primary programming language for the implementation of the project. Its powerful type system allows for a number of errors to be rejected at compile time; a quality that is helpful in developing an application with such an emphasis on fault tolerance. OCaml also provides a rich third party eco–system with libraries for crucial components of the system such as concurrency and RPCs.

In order to check the fault tolerance of the implementation it is necessary to simulate a distributed system on a simulated network. This simulator is used to check that the algorithm provides consensus in the face of a number of failure modes. It is also used to characterise the performance of the implementation in terms of latency and throughput.

Chapter 2

Preparation

This chapter describes the necessary background material required to understand Paxos; first in the simpler case of agreeing on one value and then going on to the Multi-decree case that will be implemented. From this a set of requirements for the project is established, taking into account functionality that is glossed over in theoretical descriptions of the algorithm. Finally the software engineering aspects of the project will be discussed including the choice of tools, build processes and third party libraries.

2.1 Theoretical background

2.1.1 Assumptions of the environment

When considering developing a system with distributed consensus, it is necessary to consider the assumptions made in the environment in which such a system will operate. This is to ensure there is enough functionality embedded in the conesnsus system to ensure that consensus is reached under the given set of assumptions.

A *process* (or networked process) is an instance of the program running on a networked machine in a distributed system. Assumptions of these processes that participate in the system:

- Processes can undergo *crash failures*. A crash failure is defined as a process terminating but not entering an invalid state.

- Processes may recover from crash failures. They can rejoin the system in some valid state.

- Processes operate at arbitrary speeds. This cannot be distinguished by other processes from arbitrarily long delays in the network.

Assumptions of the network in which these processes communicate:

- All processes can communicate with one another.

- The network environment is *asynchronous*. That is, messages may take an arbitrarily long time to be delivered.

- Messages may be re-ordered upon delivery.

- Messages may be duplicated in the network.

- Messages may be dropped from the network.

- Messages are not corrupted or modified in the network.

- A message that is received by one process was, at some point in the past, sent by another process.

The last two assumptions assume a system that does not tolerate what are known generally as *Byzantine failures*.

2.1.2   Aim of consensus

With distributed consensus, we wish for a network of processes to agree on some value. In consensus algorithms it is assumed that processes can somehow propose values to one participating processes. The goal of distributed consensus is, given a number of processes that can each propose some value v, that one of the proposed values is chosen. This is the *single-decree* case, that is only one value is proposed by each process and only one is chosen.

In the *multi-decree* case, agreement is reached over a sequence of values. That is, each process will propose a sequence of valus $v_1, v_2, \ldots, v_n$ and the role of the consensus protocol is to have the system choose one such sequence from all those proposed. This multi–decree case allows for the state machine replication technique to be employed to replicate an application across a number of machines in a distributed system.

2.1.3   State machine replication

A desirable goal of distributed computing is to replicate an application across a number of machines so that each *replica* has the same strongly–consistent application state. This technique is referred to as State machine replication (SMR); it leads to both increased fault tolerance and higher availability. Multi–decree consensus protocols provide a primitive by which an application (that behaves

Role	Purpose	Number required
Proposer	Propose values to acceptors. Send prepare requests with proposal numbers.	$f + 1$
Acceptor	Decide whether to *adopt* a proposal based on its proposal number Decide whether to *accept* a proposal based on a higher numbered proposal having arriving.	$2f + 1$
Learner	Learn value chosen by majority of acceptors	$f + 1$

Table 2.1: Summary of the roles in single–decree Paxos. In this description the system can tolerate the failure of up to f of each given role.

like a state machine) can be replicated.

Each process participating in the consensus protocol runs the replicated state machine application, with each process starting in the same state. By treating the values proposed in the consensus protocol as *commands* to perform a state transition, running a consensus protocol between each process cause cause each to receive the same serialized sequence of commands $c_1, c_2, \ldots, c_n$. These commands are treated as commands to perform a state transition and as such each process perform the same sequence of transitions from the same starting state and thus be a replica of the the state machine application.

Before considering how to implement SMR in the multi–decree case, it is useful to examine how Paxos operates in the simpler single–decree case.

2.1.4 Single–decree Paxos

Single–decree Paxos is the variant of the algorithm that allows for a single value to be chosen from a set of proposals and provides a foundation for the multi–decree case that will be considered next. The terminology used here follows Lamport's paper [6] describing the single–decree protocol in simple terms. Processes take the roles of *proposers*, *acceptors* and *learners*, each of which has a designated task in the algorithm. In reality these roles are often co-located within a single process but it is simpler to consider each separately. The prupose of each role and the number of each role required to tolerate f failures is summarised in Table 2.1.

Proposers that wish to propose a value v submit proposals of the form (n, v), where $n \in \mathbb{N}$ is called a proposal number. Each proposer may propose one proposal at a time and may only use strictly increasing proposal numbers for each proposal. Furthermore, each proposer must use a disjoint set of proposal numbers. The Paxos algorithm is divided into a number of stages described below.

Phase 1a (Prepare phase) A proposer wishing to propose a value first sends a `prepare(`n`)` message to a majority of the set of acceptors, where n is the highest proposal number it has used so far.

Phase 1b (Promise phase) In this phase an acceptor receiving a `prepare(`n`)` message must decide whether or not to *adopt* this proposal number. Adopting a proposal number is the act of promising not to accept a future proposal number n' such that $n' < n$. The acceptor will adopt n if it is the highest proposal number it has received thus far, in which case it will reply to the proposer with a `promise(`n''`,`v`)` message, where n'' is the highest proposal number it has previously accepted and v is the corresponding proposal's value. Otherwise, it can simply ignore the proposer or send a `NACK` message so the proposer can abandon the proposal.

Phase 2a (Accept phase) Upon receipt of a `promise(`n`,`v`)` message from a majority of the set of acceptors, the proposer replies to each with an `accept(`n'`,`v'`)`, where n' is the highest proposal number returned by the acceptors in the promise phase and v' is its corresponding value.

Phase 2b (Commit phase) An acceptor receiving a `accept(`n`,`v`)` message from a proposer will decide whether to commit the proposal for v. If the acceptor hasn't made a promise to adopt a proposal number higher than n, then it will commit v, otherwise it will ignore this message or send a `NACK` to the proposer.

Figures 2.1 and 2.2 show two example message flows between proposers and acceptors in the single Paxos protocol.

Once this process is completed, a majority of the acceptors will have chosen the same proposed value. Learners are required to each learn the value chosen by the majority. A number of different methods can be employed achieve this. The simplest is to have acceptors, on choosing a value to accept, broadcast their decision to the set of learners.

Figure 2.1: Timing diagram showing a single proposer p_1 committing a value v with three acceptors a_1, a_2 and a_3

Figure 2.2: Timing diagram showing a proposer p_1 being preempted by acceptors a_1, a_2 and a_3 having already adopted a higher proposal number submitted previously by p_2

2.1.5   Multi–decree Paxos

Single–decree Paxos can be naively extended by allowing proposers to propose values one at a time. However, this is wasteful as it requires that proposers send `prepare` messages for each proposal they wish to make. Van Renesse et al [13] describes Multi-Paxos with an emphasis on use for state machine replication that forms the basis of the description that follows. In this paper the roles of the processes are changed to emphasise the purpose of the system SMR. *Replica* processes are introduced to implicitly act as the learners and maintain the replicated state machine. *Client* processes are introduced as a means to externally submit commands to the replicas. Proposers are referred to as *leaders* as each takes the lead over some set of *ballots*. The new roles and their correspondence to the single–decree roles are summarised in Table 2.2.

Clients

The purpose of clients is to allow for commands to be sent externally to the system which can then be formed into proposals internally. This allows the system to behave in a manner more like that of a typically deployed distributed system (with a client / server architecture) and provides a degree of failure transparency.

A command c takes the form (κ, cid, op), where κ is a unique identifier for the client, cid is a unique identifier for the client's sent commands and op is the operation the command should perform. Clients broadcast a `request(`c`)` message to the replicas and each is issued a `response(`$cid, result$`)` when consensus is reached and it has been applied to each replica's application state.

Replicas

Replicas receive commands and attempt to serialize them by converting each command c into a proposal (s, c), where $s \in \mathbb{N}$ is a slot number. The slot number describes ordering of the sequence in which the commands should be committed; this is not to be confused with the proposal number n in the single-decree protocol.

Different replicas may form different sequences of proposals and so broadcasts a `propose(`s,c`)` message to the set of leaders and awaits a `decision(`s', c'`)` message. The resulting decision may differ in its slot number and so the replica may have to re-propose a command it has proposed for the decided slot. Upon receipt of

Role	Equivalent	Purpose	Number required
Client	N/A	Send commands to replicas and receive responses	*N/A*
Replica	Learner	Receive requests from clients. Serialize proposals and send to leaders. Receive decisions and apply to the replicated application state. Handle reconfiguration of set of leaders	$f + 1$
Leader	Proposer	Request acceptors adopt ballots. Secure adoption of ballots in phase 1 of synod protocol Secure acceptance of ballots in phase 2	$f + 1$
Acceptor	Acceptor	Adopt the highest ballot they have received Accept pvalues of ballot they have adopted Fault tolerant distributed memory of pvalues	$2f + 1$

Table 2.2: Summary of the roles in Multi Paxos. In this description the system can tolerate the failure of up to f of each given role. Clients do not explicitly participate in the protocol and so there is no requirement on any number being live at any given time.

decisions the replica will apply the associated operation to the application state, maintaining the replicated application.

Ballots and pvalues

Ballots are the structure over which leaders and acceptors operate when participating in the synod protocol. Ballots provide indirection in that they allow for leaders to secure agreement with acceptors for a number of proposals with different slot numbers or for a given slot to be targeted by a number of proposals associated with different ballots. This allows for a leader to perform the first phase of the synod protocol independently of having received any new proposals. It also increases the ability to execute the protocol concurrently for a number of slots at once.

Each ballot is uniquely identified by a *ballot number*. These are either pairs (r, λ) (where $r \in \mathbb{N}$ is called a round number and λ is a leader's unique identifier) or $\bot$, a specially designated least ballot number.

Pvalues associate a ballot number with a proposal. They are triples (b, s, c) consisting of a ballot number, a slot number and a command. These are analogous to to the (n, v) pairs used in the single-decree case. We previously required that each proposer used a disjoint subset of $\mathbb{N}$ for their proposal numbers. We can avoid this requirement as each ballot number encodes the identifier of the leader directly in its ballot number. Hence no two leaders can generate equal ballot numbers.

We require ballot numbers to be totally ordered so that acceptors can compare which ballot number is less than another when choosing whether to adopt or accept. Letting $\mathcal{B}$ denote the set of all ballot numbers, we define the relation $\leq_{\mathcal{B}} \in \mathcal{B} \times \mathcal{B}$ which satisfies the following two conditions:

$$\forall (n, \lambda), (n', \lambda') \in \mathcal{B}. \ (n, \lambda) \leq_{\mathcal{B}} (n', \lambda') \iff (n \leq n') \vee (n = n' \wedge \lambda \leq_{\Lambda} \lambda')$$
$$\forall \, b \in \mathcal{B} \, . \, \bot \leq_{\mathcal{B}} b$$

Note this implies that we require leader identifiers be equipped with a total order relation $\leq_{\Lambda}$ as well.

Quorums

Unlike replicas, acceptors do not each attempt to store their own consistent copies of proposals. Instead, it is required that at least a majority of the acceptors have in their memory a decision for a committed proposal. This is so that if any minority of the acceptors crash then there is still going to be at least one acceptor with memory of the committed proposal. A subset of the acceptors $Q \subseteq \mathcal{A}$, where $|\mathcal{A}| = 2f + 1$ for some $f \in \mathbb{N}$, is called a *majority quorum* or just *quorum* if $|Q| = f + 1$.

Hence at each phase of the synod protocol leaders broadcast their message to all acceptors in a configuration and wait for such a quorum of responses. This is to guarantee that if any minority of the acceptors fail in the future that at least one will have retained the result of the current round of the protocol.

The Synod Protocol

The synod protocol is the protocol undertaken by the set of leaders and acceptors in order to decide which command is committed to which slot. The protocol proceeds in two phases similarly to the single-decree case except now leaders and acceptors operate over pvalues.

Acceptors maintain a ballot number in their state their represents the last ballot number they have *adopted*. Adopting a ballot number represents a promise not to accept any pvalues that do not have that ballot number, so that acceptors do not accept conflicting pvalues. Acceptors *accept* a pvalue by adding it to the set of pvalues, initially empty, that they have already accepted. The algorithm requires that if two acceptors $\alpha, \alpha' \in \mathcal{A}$ have accepted (b, s, c) and (b, s, c') then $c = c'$.

In the absence of receiving any `propose(s,c)` messages from replicas, leaders attempt to secure an initial ballot with ballot number $(0, \lambda_k)$, where λ_k is the identifier of the k^{th} leader. They do this by broadcasting a `phase1a(b)` message in the same format as that below.

Phase 1a Leaders attempt to secure adoption of a ballot b by broadcasting a `phase1a(b)` message to the set of acceptors.

Phase 1b Acceptors receiving a `phase1a(b)` compare b to the highest ballot number b' they have adopted thus far. This is initially $\bot$ so acceptors will always

adopt the first ballot number they receive. If $b' \leq_{\mathcal{B}} b$, then the acceptor will adopt this new ballot number, that is set $b' := b$. In either case, the acceptor will reply with a `phase1b`(b', `pvals`) message, where pvals is the set of pvalues previously accepted.

If a leader receives `phase1b`(b', `pvals`) from a majority quorum of acceptors and if $b = b'$ then the leader will proceed with phase 2 of the protocol. Otherwise the leader abandons the proposal, increments the round number of its ballot number and tries again. If a leader receives a response where $b' \neq b$ then it should restart the protocol with a higher round number.

Phase 2a Leaders proceed to commit a given pvalue (b, s, c) by broadcasting a `phase2a`(b, s, c) message to the set of acceptors.

Phase 2b Acceptors receiving a `phase2a`(b, s, c) will compare b to their adopted ballot number b'. If $b = b'$ then they will accept the pvalue (b, s, c); otherwise they do not change their state. In either case they reply with a `phase2b`(b') message.

Leaders will again wait for receipt of a `phase2b`(b') from a majority quorum of acceptors. If $b = b'$ then they know the acceptor has accepted the pvalue and so broadcast a `decision`(s, c) message to the set of replicas so they can commit command c to slot s. Otherwise if $b \neq b'$ then the acceptor knows the acceptor has since adopted a different leader's ballot number and so should restart the protocol with a higher round number.

2.2 Requirements

2.2.1 System requirements

As noted previously, there are a number of papers that describe the Paxos protocol and its variants. Many of these sources present the algorithm in a different format and nearly all of them present it in a theoretical setting, with little concern for functionality that is generally required in software. Implementation details such as how each of the participating process knows how to address one another are often entirely overlooked. Hence, in developing an implementation of Multi-Paxos, it is necessary to formally recognise the theoretical requirements

from the literature as well as identify the additional functionality required to realise the algorithm as a functioning piece of software.

It is therefore necessary to perform a *requirements analysis* of the final system. This was performed by collecting information from the literature on how the algorithm is presented and also considering all of the necessary functionality required to implement such an algorithm. For example, in the Paxos Made Moderately Complex paper processes are described as being able to send messages between one another. This of course presents the requirement for a messaging primitive that operates on top of a best–effort IP network. The final system requirements are listed below.

General requirements

- The ability to pass the executable file some command line arguments for initializing the process.

- Ability for each node to be nominated to a given role and to start-up in that role (via the command line).

- Configuration files that contain the network addresses of each of the processes participating. Each process must be able to parse such a configuration file.

- The ability to log important information to disk or direct to `stdout` / `stderr`. The ability to log different information (e.g. debugging info versus error info) to different files.

Messaging subsystem

- Provide an abstraction over the network so processes can send messages, rather than worrying about functionality of network.

- Send messages as required by the Multi-Paxos algorithm, being able to serialise and deserialise arguments and results.

- Mask failures in the network from processes

Application requirements

- A key value store application that can be modelled as a state machine. This needs to have commands to create key value pairs, update key values pairs, read a value for a given key and delete a key value pair for a given key.

- The store must be strongly consistent across all replicas.

Consensus algorithm requirements

- Clients processes can messages with commands to the replicas and receive responses.

- Replicas that each maintain a strongly consistent copy of the application state. They must receive requests from clients and each attempt to propose their own serialisation of the sequence of the requests to leaders.

- Leaders and acceptors take part in the Synod protocol that produces a consistent serialisation of the commands proposed by replicas.

2.3   Software engineering

This section gives detail on the software development pratices and related work that was undertaken to let development commence. The development methods, third party libraries, tools and test strategies that were employed to manage the project and improve productivity are discussed here.

2.3.1   Starting point

The starting point of this project is the same as that described in the Project Proposal reproduced in Appendix B.

2.3.2   Methodology

The method by which the software was developed was to split the system into its logical constituents and develop each in turn and with as much independence from the rest as possible. To facilitate this and from the requirements I deemed it necessary to separate the messaging system from the actual operation of Multi–Paxos itself.

This messaging system was the first module to be developed and was incrementally equipped with functionality that facilitated the rest of the program. The next logical division was to separate the operation of clients and replicas from that of leaders and acceptors. Clients, replicas and associated functionality was developed first followed by that of leaders and acceptors in two phases of development.

2.3.3 Libraries

OCaml has a large eco-system of third party libraries available. They are used where it is necessary to avoid writing large bodies of code for common tasks. This project uses the OPAM[1] package manager to install, update and manage any third party libraries via a command line interface.

A library of particular interest to this project is that which provides RPC (remote procedure call) functionality. Cap'n Proto[2] was chosen as it has good support and documentation for its OCaml bindings; separated into Capnp-ocaml[3] for serialization and Capnp-rpc[4] for RPCs itself. The library provides a schema language to describe the required remote procedures and any required structures and a code generator for generating OCaml bindings.

Since this project is in the space of distributed systems it is necessary to consider how concurrent programming would be achieved. Lwt[5] is a monadic concurrency library that leverages the type system to express the promise of a computation, of type `'a Lwt.t`, that returns a value of type `'a`. Monadic functions `return` and `bind` to chain computations together in a safe manner and support for parallel computation, mutexes and Unix bindings. This library was chosen because it has thorough documentation and interoperates with Cap'n Proto.

Whilst OCaml already comes with a standard library, this project uses Core[6] as a popular overlay. Core improves upon the standard library by providing a number of additional useful functions and modules and also in some places improves the type definitions of those present in the default standard library (by introducing named arguments in order to improve readability). Yojson[7] was used for serialization, both in configuration files and in some areas sending data in messages. OUnit[8] was used to manage unit testing for this project and is discussed in 2.3.6.

In the case of each library used their software lisence was made readily available

---

[1]https://opam.ocaml.org
[2]https://capnproto.org/
[3]https://github.com/capnproto/capnp-ocaml
[4]https://github.com/mirage/capnp-rpc
[5]https://github.com/ocsigen/lwt
[6]https://github.com/janestreet/core
[7]https://github.com/mjambon/yojson
[8]http://ounit.forge.ocamlcore.org/

Library	License	Information availabble
Lwt	LGPL with OpenSSL linking exception	`opam show lwt`
Core	Apache-2.0	`opam show core`
YoJson	Specific license file	`https://github.com/mjambon/` `yojson/blob/master/LICENSE.md`
OUnit	Specific license file	`https://github.com/mjambon/` `yojson/blob/master/LICENSE.md`
Ocaml-capnp Capnp-rpc	Apache	`opam show ocaml-capnp` `opam show capnp-rpc`

Table 2.3: Software licenses for each of the libraries used in the project.

either by OPAM (via the `opam show <pkg-name>` command) or by licensing information made available on a library's Github repository. Each of the libraries used had licenses that made use of the software for this project possible, with YoJson and Uri requiring reproduction of their licensing information distributed with the project. Table 2.3 gives details of each library's individual license.

2.3.4 Build system

JBuilder was used as the build system for this project. It was chosen for its minimal configuration files (called jbuild files), automatic linking of OPAM packages and automates the generation of `.merlin` files. Jbuild files are written in a simple S-expression syntax and are composable with one another. Jbuild files also present the ability to include *rules* so that Cap'n Proto OCaml bindings are generated automatically when building the project.

2.3.5 Tools

Choice of tools is important for being productive in a language. MacVim[9] was used as a text editor. In order to improve productivity Merlin was used to provide code-completion and edit-time type inference features, making it much easier to understand the types of source code and save time reading documentation online. `.merlin` files are auto-generated by Jbuilder that allow for code completion and type inference over files in the project directory itself.

---

[9]https://github.com/macvim-dev/macvim

Git was used for source control and version. Branches were made for each new feature under development and merged back into the master. Backups were to a remote repository on Github. Backups were also made periodically to Google Drive. All of the development work was undertaken on my person machine with the option of fallback onto the MCS machines.

In order to simulate the operation and performance of instances of the implementation in a distributed system it was necessary to use a network simulator. Mininet[10] is a network simulator that operates with proess based virtualisation and allows for arbitrary processes to be simulated as running on hosts on a virtual network. Mininet provides a complete Python API for setting up such simulations and performing experiments. Mininet required use of a guest operating system on VirtualBox[11], virtualization software that allowed Mininet to be used on my own machine.

2.3.6 Testing

Unit testing was performed in order to check the operation of functions at a fine-grained level within the program. OCaml's static type system ensures the compiler will catch a number of errors resulting from improper calls to functions and so unit tests were designed as a compliment to this; tests focussed on correctly typed arguments to functions yielding appropriate results we expect. The test suites were executed after successful builds of the project.

2.4 Summary

This chapter discussed the theoretical basis for implementing Multi–Paxos and state machine replication before going on to describe a set of requirements for an implementation of the algorithm, finishing with a description of the software engineering tools and techniques employed. With this in place it is possible to go on to describe the implementation of the algorithm and its underlying systems in Chapter 3.

---

[10]http://mininet.org/
[11]https://www.virtualbox.org/

Chapter 3

Implementation

This chapter describes the implementation. It begins with a high level overview of the structure of the program and initialisation. This follows with an examination of the data structures that are used and goes on to describe the messaging system. Finally, each of the roles in the system is described in turn.

3.1 High level architecture

This section describes the architecture of the program. Each process taking participating in Multi-Paxos is an instance of the same executable configured to take the *role* of either client, replica, leader or acceptor. These roles each have a different function in the system within their *role module*, which ecapsulates all the functionality associated specifically with a given role. Each of those modules will need to send and receive messages of some form; this is abstracted away by a *messaging system* which exposes an interface via which each can send messages or maintain a server that receives messages. Figure 3.1 shows communication paths between each of these subsystems.

Each of the processes requires additional information when they are initialised, such as

1. The role it should take.

2. The address at which it should start a server to receive messages.

3. The set of addresses of all the participating processes with which it must communicate.

Items (1) and (2) are provided as comand line arguments when the program is executed. If the number of processes is large, then supplying a list of addresses to each process would prove tedious so (3) is addressed by using a configuration file, the relative path to which is provided as a command line argument.
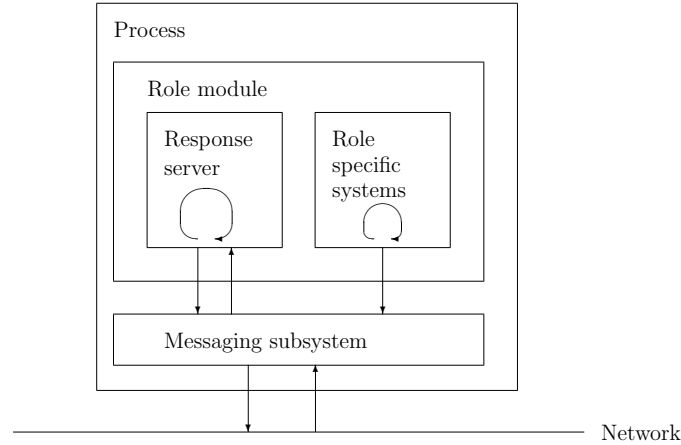
Figure 3.1: High level subsystems of a process

Configuration files are formatted as JSON[1]; JSON was chosen since it is semi-structured data that can be easily handwritten and also parsed using the Yojson library. The file contains, for each role, a list of all of the addresses of each process in that role. Not every process needs to know the address of every other node. For example, clients do not need to know the addresses of leaders and acceptors. It is only necessary that each file parsed by each process contains the addresses of the nodes they are required to communicate with. In practice, it is more consistent and maintainable to have a single file used by each of the process.

The addresses discussed above are IPV4 IP address and port number pairs which are used both to send and receive. Internally a different addressing format is used, described in 3.3, where the addresses provided at initialisation are used to map directly onto the internal representation.

3.2 Data structures

Before proceeding the implementation of the systems described above, it is necessary to examine the data structures that will be used.

3.2.1 Identifiers

Unique identifiers are required to identify each process in the system. To avoid having a central authority distribute these identifiers a method where any

---

[1]https://tools.ietf.org/html/rfc7159

```
module type APPLICATION = sig
  type state = (int, string) List.Assoc.t
  type operation = Nop
                 | Create of int * string
                    ...
  type result = Success
              | Failure
                 ...
  val initial_state : state
  val apply : state -> operation -> state * result
   ...
end
```

Figure 3.2: Signature of the key value store module

process can generate their own unique identifer is used. Universally Unique
Identifers[2] (UUIDs) are used in this case as they are a well established standard
and have support in the OCaml Core library. UUIDs are also totally ordered
and so satisy the condition that we have such an ordering on the identifers of
leaders.

Hence replicas, leaders and acceptor identifiers have the type `Core.Uuid.t`. As
we wish to expose clients to a request / response protocol it is necessary to store
a map of client identifiers to addresses by each replica. Rather, by modifying
client identifiers to be of the form `Core.Uuid.t * Uri.t`, the address actually
forms part of the identifier.

3.2.2   Key value store

The key value store is the application that is to be replicated. Each replica
process will maintain its own independent copy of its state. The application need
not maintain any synchronisation logic, it need only behave as a state machine.
The state is represented by an association list that maps integer keys to string
values. Operations each have their own semantics described in Table 3.1. The
application follows a state machine pattern given that if each replica starts in
the initial state and each applies the same sequence of operations in the same
order, the resulting state is the same.

---

[2]https://tools.ietf.org/html/rfc4122

Command	Argument	Semantics
Nop		No operation, no change to state Returns Success
Create	(K,V)	If key not present, add new (K,V) pair to the state and return `Success`. Otherwise return `Failure`.
Update	(K,V)	If key present, update pair with key K to value V. Otherwise return `Failure`
Read	K	If key present read value V associated with key K and return `ReadSuccess(V)`, else return `Failure`.
Remove	K	If key present remove pair with key K, else return `Failure`.

Table 3.1: Set of operations that can be applied to the application state, along with their corresponding arguments and their semantics. Note that Create and Update commands are separate, each with their own success and failure semantics in order to reduce ambiguity of how the application operates.

```
type t = Bottom
       | Number of int * leader_id
```

Figure 3.3: Types of ballot numbers.

When a command has been committed to a slot by the consensus algorithm, the command's operation is applied to the state. This returns a new state and a result. The result is the value that is returned to the client; either `Success`, `Failure` or `ReadSuccess(V)`. The signature of the key value store is shown in Figure 3.2.

3.2.3 Ballots

The definition of ballot numbers from 2.1.5 lends itself to representation by an algebraic datatype. Ballots are hence the tagged union of Bottom (representing the least ballot $\perp$) or a pair consiting of an integer and a leader identifer (representing pairs (n, λ)). The type definition is listed in 3.3

Figure 3.4 shows the interface exposed by the Ballot module. Note that the concrete type of a ballot number isn't exposed, only the abstract type `t`.

```
type t
val bottom : unit -> t
val init : leader_id -> t
val succ_exn : t -> t
```

Figure 3.4: The interface exposed by the Ballot module. These are the types of functions that can be used to generate ballot numbers. Note the naming of the function succ_exn implies it can throw an exception, which occurs when calling the function on bottom.

Therefore outside of this module the internal representation of ballot numbers is hidden. This prevents a large number of errors, such as decrementing a round number when they should strictly increase, from arising by rejecting them at compile-time. It is necessary to supply function to work with ballot numbers outside the module.

Hence a given leader can generate an initial ballot number with their identifer. Subsequent ballot numbers can be generated by calling a successor function, which increments the round number each time. By exposing such an interface we prevent errors such as using negative round numbers from being able to compile.

The module provides functions to test equality and the ordering of ballots. These work over the structure of ballots, comparing them first for correct types and then checking round number and leader id equality. Also in the Ballot module are functions for serialisation and deserialisation to and from JSON, required by the messaging system for sending ballots.

3.3   Messages

3.3.1   Interface exposed

The messaging subsystem is concerned with the sending and receipt of messages between processes. It transforms messages from simple OCaml types into a form suitable for transport over an IP network, handling packetisation, rentramssion and masking failures. It also handles the initialisation of servers and the addressing in the system.

There are a number of different messages that are necessary to send in Multi-Paxos. In this chapter we focus on the capability of sending these messages

```
type non_blocking_message = ClientRequestMessage of command
                          | ProposalMessage of proposal
                          | DecisionMessage of proposal
                          | ClientResponseMessage of command_id *
                            result

val send_non_blocking : non_blocking_message -> Uri.t -> unit Lwt.t

val send_phase1_message : Ballot.t -> Uri.t ->
  (Core.Uuid.t * Ballot.t * Pval.t list, string) Result.result Lwt.t

val send_phase2_message : Pval.t -> Uri.t ->
  (Core.Uuid.t * Ballot.t, string) Result.result Lwt.t
```

Figure 3.5: Types of non-blocking messages

```
val start_new_server : ?request_callback:(Types.command -> unit) ->
    ?proposal_callback:(Types.proposal -> unit) ->
    ?response_callback:(Types.command_id * Types.result -> unit) ->
    ?phase1_callback:(Ballot.t -> Types.unique_id * Ballot.t * Pval.
        t list) ->
    ?phase2_callback:(Pval.t -> Types.unique_id * Ballot.t) ->
    string -> int -> Uri.t Lwt.t
```

Figure 3.6: Function to start a new server that presents a number of possible callbacks

rather than their role in the algorithm itself, which is discussed in depth over the rest of this chapter.

Messages are broadly divided into two cateogries: *blocking* and *non-blocking*. Blocking messages are those which expect a direct response to a given message, whereas when sending a non-blocking message the process does not expect an immediate reply. The interface for sending messages is presented in Figure 3.5. The types of non-blocking messages represent the arguments that are associated with such a message and do not include any information about source and destination addressing, as this is handled transparently by the messaging system.

The function send_non_blocking takes a non-blocking message represented by the type described above and a URI that addresses the destination process. As these messages do not explicity require a reply they return a unit Lwt.t, that is a promise to perform a computation that terminates with value (). As we can

regard a message that is never delivered to a recipient as indefinitely delayed in the network, we return this type regardless of any errors encountered by the underlying subsystem. It is important to note here that any errors that arise from messages failing to deliver are handled by the Multi-Paxos algorithm itself, not the messaging system.

Blocking messages on the other hand return a type immediately. Hence they are separated into their own specific functions, `send_phase1_message` and `send_phase2_message`, each with a return type that represents the arguments in the response. Note that here the arguments represent the phase1a / phase2a messages and the responses represents the phase2a and phase2b messages. We do this because acceptors only respond in response to a message from a leader, so it reduces the functionality required to represent these replies as return types. The error handling here returns a type wrapped in a `Result.result` type which allows for an error to be returned and discarded by the leader when checking whether it has received responses from a quorum.

The message module also exposes a function to start a new server, the type of which is given in 3.6. The server requires a string and integer representing the IP address and port number on which to listen and a number of named functions that each represents a callback. These functions are called when a corresponding message is received by the server and return the response given by the return type. A process can therefore interoperate with the server by passing functions it wishes to have called whenever a message is received.

3.3.2   Cap'n Proto

Cap'n Proto is used as the underlying system for sending messages. Cap'n Protos requires that one writes a schema file that describes a service; this contains data structures and messages that need to be serialized. This schema file, listed in full in Appendix A, contains an interface that describes the service that we wish to make available to the messaging system. The schema crucially contains a method for every message that can be sent and its associated arguments and return type, with each mapping to the associated message type described in Figure 3.5. Note that some of the arguments in the Cap'n Proto schema file are stored as `Text` types rather than as the types in 3.5 as these are first serialized into JSON using Yojson and then converted to strings. This was to provide a degree of separation in the source code. For example, this helps

```
let service_from_uri uri =
  try Lwt.return_some (Hashtbl.find sturdy_refs uri)
  with Not_found ->
    let client_vat = Capnp_rpc_unix.client_only_vat () in
    let sr = Capnp_rpc_unix.Vat.import_exn client_vat uri in
      Sturdy_ref.connect sr >>= function
        | Ok capability ->
            Hashtbl.add sturdy_refs uri capability;
            Lwt.return_some capability
        | Error _ ->
            Lwt.return_none
```

Figure 3.7: Function that returns a Cap'n Proto capability for the message service of a given URI

avoid situations where a new command for the replicated application would require rewriting of boilerplate code for messaging and recompilation of the schema.

A compiler tool is used to generate OCaml code that contains signatures and structures for the message API defined in the schema. The function start_new_server returns a service object that has a method for handling each of the possible functions that represent a message receipt for that server. In each of these functions the arguments are deserialised from the Cap'n Proto representation into the message type used by the system and the callback function called. The result of calling this function is then serialised and sent as a message to the sending process.

Sending a message requires translating a Cap'n Proto URI[3] that represents the server into a *capability* which represents a stateful connection to the server. For each server initialized with the start_new_server function the supplied IP address and port number are mapped to a URI that is used to address the internal capability. Given a URI one can derive a sturdy reference to the capability and connect via that sturdy reference. Each connection runs over TCP and so rather than have each message setup and teardown a TCP session a cache of connected capabilities is stored by the messaging module that maps URIs to capabilities, listed in Figure 3.7. The function searches the table hashed by URIs first to see if a connection has already been established and returns that if it does. Otherwise if there is a cache miss an attempt to connect to the capability is initiated. If this results in a

---

[3]https://tools.ietf.org/html/rfc3986

```
type t = {
  id : replica_id;
  mutable app_state : app_state;
  mutable slot_in : slot_number;
  mutable slot_out : slot_number;
  mutable requests : command list;
  mutable proposals : proposal list;
  mutable decisions : proposal list;
  mutable leaders : Uri.t list;
}
```

Figure 3.8: Types of records representing replica state

success then the capability cached and then returned. Note here the return type is `Capability.t option Lwt.t` since the connection could result in an error.

3.4   Clients and replicas

3.4.1   Clients

Clients are the simplest processes in the system as they need only send commands $(\kappa, \mathrm{cid}, \mathrm{op})$ to replicas in $\mathrm{request}(\kappa, \mathrm{cid}, \mathrm{op})$ messages, where κ is an identifier of the client of type `Uuid.t * Uri.t`. The identifier value is used to uniquely identify the client and the URI is a Cap'n Proto address required for replicas to direct their responses to clients in question.

Each client also maintains a server on a given host name and port number to which $\mathrm{response}(\mathrm{cid}, \mathrm{result})$ messages are sent. This represents the result of command with identifier cid having been applied to the application state, its order with all other commands having been serialized consistently across the system by Multi-Paxos.

3.4.2   Replicas

Having discussed the functionality of clients it is necessary now to talk about the implementation of replicas. The state of a replica is stored as a record with mutable fields as presented in Figure 3.8. Replicas are required to implicitly take the role of learners in order to learn the decisions made on how to serialize the sequence of client requests. Rather than have the replicas explicity request the

result of the synod protocol from the leaders or acceptors, acceptors forward the requests made by clients onto them and receive a response, implicitly learning the result.

Replicas maintain three queues:

- `requests:` A queue of commands received in request messages from clients.

- `proposals:` A queue of proposals. These are commands that have been tagged with a provisional slot number. This queue forms the replica's own attempt at serializing the commands. However, these are not committed in this order until the configuration of leaders and acceptors has decided each command for that slot.

- `decisions:` A queue of decisions. This represents the serialization of commands that has been decided on by the configuration of leaders and acceptors; it represents the consensus reached on the sequence of committed commands. It is in this order that the commands are applied to the replicated application state.

These queues are stored by the replica as types `command list` and `proposal list`. This is because list types in OCaml have a large library of helper functions in the Core library. Further, the type `proposal list` is structurally equivalent to `(slot_number, command) List.Assoc.t`, that is an association list of commands keyed by their slot number. This allows for proposals to be looked up by their key or inverted and looked up by their command, functionality necessary for managing the flow of proposals through the queues.

In order to track the next slots in which to propose and commit, each replica maintains two counters:

- `slot_in`: Represents the lowest next available slot for which no proposal has yet been submitted; that is $\nexists\ (s, c) \in$ proposals. $s = $ slot_in.

- `slot_out`: Represents the lowest next available slot for which no decision has been committed; that is $\nexists\ (s, c) \in$ decisions. $s = $ slot_out.

Commands and proposals move through this system of queues as shown in Figure 3.9. `request(`c`)` messages sent by clients arrive at the replica and are entered into the requests queue. Concurrently the replica will attempt to dequeue commands in producer–consumer relationship, proposing each to `slot_in`, the lowest free slot, and incrementing `slot_in`. Each of these

commands is tagged with this slot number and entered into the proposals queue. Once in this queue, it can then be proposed to the configuration participating in the synod protocol by broadcasting a `propose(`s`,`c`)` to the set of leaders.

Leaders will at some point in the future return a `decision(`s`,`c'`)` message for each slot s for which a command was proposed by a replica. If $c = c'$ for a proposal $(s, c) \in$ proposals then it is removed from proposals and added to decisions. However if $c \neq c'$ then a different command has been decided for the slot than the one proposed by this replica. (i.e. the replicas disagreed on their serializations and the other has won out). In this case (s, c') is committed to decisions and c is returned to requests so that it can be re-proposed for a different slot number.

Every time that a proposal is committed in the decisions queue, the proposal $($`slot_out`$, (\kappa, \text{cid}, \text{op}))$ has its command executed, followed by `slot_out` being incremented. Even though the decisions will be committed in the same sequence by each replica, the order in which they are committed may differ. By sweeping through the decisions queue and executing each we ensure the application state is updated in the same order by each replica. The application state is updated by applying op to the replica's application state, resulting in an updated state and a result res. A `response(cid,res)` message is sent to the client with identifer κ^4.

Reconfigurations represent a change in the set of leaders and acceptors participating in the synod protocol. Replicas are therefore expected to be able to reconfigure, that is to change the set of leaders with which they share their proposals. Reconfigurations are commands sent like any other, except that when executed by the replica (having been decided a slot in the same fashion any other command would), rather than being applied to the application state they update the set of leaders to the new configuration.

A window of commands is maintained so that only a maximum number of commands are ever being decided upon by the synod protocol at any given time; reconfiguration commands do not take effect until there are no more commands in flight for the previous configuration. We therefore maintain a `WINDOW` value that represents the number of slots to which a value can be proposed before a given reconfiguration command takes effect. Equivalently this involves

---

[4]Recall client identifiers have a URI component so responses can be sent without replicas maintaining a mapping of identifiers to URIs.

Clients submit requests via request(c) messages

c_4

c_3

c_2

c_1

Requests

Re-submit proposals that
conflict with decisions

(s_2, c_2)

(s_1, c_1)

Proposals

Proposals broadcast to leaders

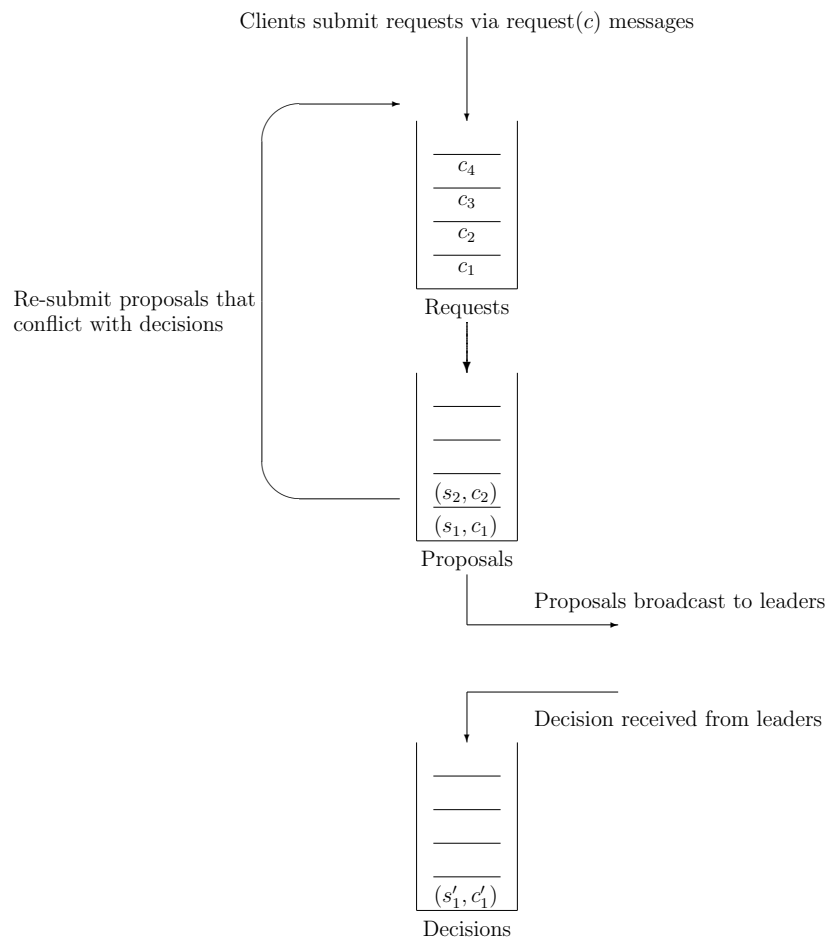Decision received from leaders

(s_1', c_1')

Decisions

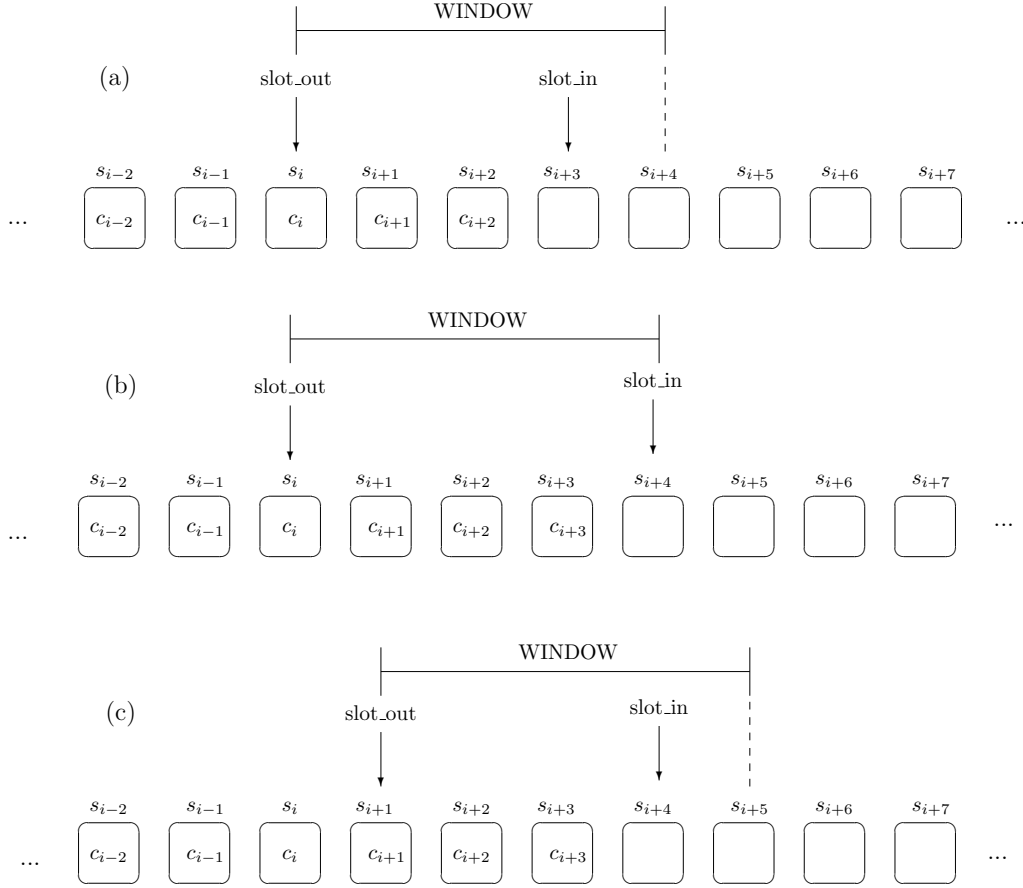Figure 3.9: Flow of requests and proposals through a replica's queues

Figure 3.10: (a) shows a window with one available slot into which a command c_{i+3} is proposed, leading to the window shown in (b). In (b) however no further slots are available for proposals since slot_in < slot_out + WINDOW. IN (c), a decision is received for c_{i+1} and so the slot_out is incremented and now another slot is available for proposals.

maintaining the invariant slot_in < slot_out + WINDOW. Hence the queue of proposals has a maximum capacity of WINDOW and implies that a maximum WINDOW proposals are in flight (that is, being serialized by the synod protocol) at once. This therefore means that given a reconfiguration doesn't take effect for WINDOW slots that after any reconfiguration is triggered all pending proposals will have been decided by the time the configuration has taken place. A diagram showing a sequence of states an example window undergoes is shown in 3.10

```
type t = {
  id : unique_id;
  mutable ballot_num : Ballot.t;
  mutable accepted : Pval.t list
}
```

Figure 3.11: Types of records storing acceptor state

3.5 Synod protocol

Replicas each propose their own serialization of client commands; varying processing speeds and delays in the network along with crash failures of replicas can cause these serializations to differ. In this case, it is the function of the Synod protocol to produce a single serialization of those proposals that is returned to all the replicas. This serialization must be decided upon in a distributed manner by the configuration, that is the set of leaders and acceptors. Before we examine the role that leaders play in the next step of the algorithm it is necessary to observe the functionality of acceptors.

3.5.1 Acceptors

Acceptors form what is known as the disitributed fault tolerant memory of the consensus protocol. As with the other processes, the internal state is stored in a record with type given in Figure 3.11. Acceptors maintain the following mutable state

- ballot_num: The ballot number that has been most recently *adopted* by the acceptor. Initially, ballot_num is equal to $\perp$ so that an acceptor adopts the first ballot it receives (following from $\perp$ being the least ballot).

- accepted: The list of pvalues which the acceptor has *accepted*, initially empty.

Acceptors represent passive processes in the algorithm. They change their state and send messages only in direct response to messages from leaders; other than when initialising, all processing occurs upon receiving a phase 1a or phase 2a message from a leader. Given this, all an acceptor is required to do is maintain its state and provide two callbacks, listed in Figure 3.12 and 3.13, when starting a server; one to receive phase 1a messages and one to receive phase2a messages. Recall from 3.3 that these are blocking messages and so the phase 1b and 2b

```
let phase1_callback (a : t) (b : Ballot.t) =
  Mutex.critical_section callback_mutex ~f:(fun () ->
    if a.ballot_num < b then
      a.ballot_num <- b);
    (a.id, a.ballot_num, a.accepted))
```

Figure 3.12: Function called when an acceptor receives a phase1a message with associated ballot.

```
let phase2_callback (a : t) (pval : Pval.t) =
  Mutex.critical_section callback_mutex ~f:(fun () ->
    let (b,_,_) = pval in
    if b = a.ballot_num then
      (if not (List.mem a.accepted pval ~equal:(Pval.equal)) then
        a.accepted <- pval :: a.accepted);
    (a.id, a.ballot_num))
```

Figure 3.13: Function called when an acceptor receives a phase2a message with associated pvalue.

replies are just the return type of these callbacks.

Both the phase 1 and phase 2 callback functions can be executed concurrently as each callback is executed when the underlying server receives a message. Hence it is necessary to enclose both the callbacks within critical sections to ensure data mutated in one callback isn't used in another. The function Mutex.critical_section takes a function that contains the contents of the callbacks and the mutex to lock when entering the critical section. Hence when one function holds the lock any other callbacks will block upon attempting to hold the lock and receive the lock when the function holding the lock exits the criticial section.

Leaders broadcast phase1_callback(b) to the set of acceptors when they seek an adoption of ballot number b form a majority quorum of acceptors. In this case, when an acceptor receives a phase1_callback(b) message, it will adopt the ballot if it is greater than the last ballot it adopted (and hence inductively it is greater than any previous ballots it has ever adopted). The acceptor will then return a triple consisting of its identifier, the ballot it has adopted (this will either be the ballot it has just received or the ballot it previously adopted) and the list of accepted pvalues. The identifier of the acceptor is returned so leaders

can track which acceptors have replied. Rather than just not reply when it does not adopt b, it responds with the already adopted ballot so that leaders can be *preempted* and abandon this adoption attempt and try a higher ballot number. The list of accepted pvalues is returned so that if the acceptor has accepted any pvalues from other leaders in the past the leader attempting to secure adoption of b will learn about them from this reply.

Leaders send `phase2_callback(pval)` messages, where `pval` $= (b, s, c)$, when they believe b has been adopted by a majority quorum of acceptors; this is the leaders request to commit the proposal (s, c). When an acceptor receives such a message they may have adopted a higher ballot number or the message may even have been delayed. Hence if the acceptor's `ballot_num` $= b$ then the acceptor does accept the pvalue (b, s, c) and adds it to the list of accepted pvalues; otherwise it doesn't. In either case, it returns its identifier and `ballot_num`, so that in case the leader will know whether (b, s, c) was accepted or a higher ballot had been adopted by the acceptor.

In this situation each acceptor may have a different list of accepted pvalues. When we note that leaders only decide that a proposal (s, c) is committed after receiving a majority quorum of phase 2 messages then each such proposal must have a corresponding ballot number (b, s, c) accepted a majority of acceptors. This is what gives us the fault tolerant memory, we can tolerate the crashes of a minority of acceptors (f out of $2f + 1$ total participating) and still have a consistent memory of the sequence of proposals.

3.5.2   Leaders

Leaders receive over time from each replica their serialization of the set of proposals which may differ. The leaders are required to, between them and in a manner that tolerates the failure of all but one leader, return a serialization of the proposals such that each replica receives the same serialization.

Leaders spawn and manage *scouts* and *commanders* in order to separate the concerns of attempting to secure adoption of ballots and attempting to secure acceptance of pvalues. Each scout and commander is a sub-process with an associated state and execution context that is spawned by a given leader. Rather than using the existing messaing subsystem which would produce excessive overhead for managing exchange of data between these threads we instead have

```
type process_response = Adopted of Ballot.t * Pval.t list
                      | Preempted of Ballot.t
```

Figure 3.14: Types of responses produced by scouts and commanders

```
let queue_guard = Lwt_mutex.create ()
let message_queue : process_response Queue.t = Queue.of_list []

let send msg =
  Lwt_mutex.lock queue_guard >|= (fun () ->
  Queue.enqueue message_queue msg) >|= fun () ->
  Lwt_mutex.unlock queue_guard
```

Figure 3.15: Functions and values for manipulating the message queue of process responses

them communicate locally in a producer consumer relationship.

In essence each sub-process has a lifetime in which it is spawned, messages acceptors, receives responses, performs some processing and then terminates. The terminating sub-process may have to pass back to the leader the result of its computation. In the case of scouts, this is either a notification of a preemption having occurred or its ballot having been adopted (along with the pvalues of all the previously accepted ballots). Commanders may notify the leader on termination of a preemption having occured. The types of these responses are listed in Figure 3.14.

We treat the sub-processes and the leader that spawned them as engaging in a producer consumer relationship. Figure 3.15 lists code pertaining to the response queue. Terminating sub-processes produce these responses and enqueue them onto message_queue, a queue of responses, via the send : process_response -> unit Lwt.t function. The leader periodically will consume these responses from the queue, in turn possibly spawning more sub-processes that repeat the cycle. The queue is protected by a guard mutex queue_guard that protects from races from occuring when the leader and one or more sub-processes attempt to mutate the state of the queue at once.

```
module type SCOUT = sig
  type t' = {
    b : Ballot.t;
    acceptor_uris : Uri.t list;
    receive_lock : Lwt_mutex.t;
    mutable pvalues : Pval.t list;
    mutable quorum : (Uri.t, unique_id) Quorum.t
  }

  val spawn : t -> Ballot.t -> unit
end
```

Figure 3.16: Signature of scouts

Scouts

As described above, the purpose of a scout sub-process is to take a given ballot number and attempt to secure adoption of that ballot number with a majority quorum of acceptors. The ballot is spawned by a commander, passed this ballot number and terminates by enqueuing a Adopted(b,pvals) or Preempted(b') response for the leader to consume.

Scouts are described by a signature listed in 3.16. Each scout's state is represented by a mutable record of type t'. Calling spawn and passing as arguments a leader and a ballot number will start a initialize a new scout and start its asynchronous execution. Since its execution is asynchronous the function returns () immediately. Note that the ballot number b over which it seeks adoptions is immutable; even though a leader works over ever increasing ballot numbers, b is fixed for the duration of the a given scout's lifetime.

Upon being spawned, a scout broadcasts (in parallel) to the set of acceptors a phase1a(b) message. Each of these is bound monadically with a function that returns unit if an error occured in message delivery (in keeping with our messaging semantics) or the phase 1a return value $(\alpha, b', \mathsf{pvalues}')$. Upon receiving each message, the scout checks if it has already obtained a majority quorum of responses (not including the response just received), in which case it discards the response.

If $b = b'$ then the scout adds pvalues' to the pvalues it already has stored and adds α to the quorum of responses it has received. If the scout has now obtained a majority quorum then the scout terminates, adding a Adopted(b,pvals)

message to the queue for a leader to consume. If not, the scout continues to wait for further responses from the acceptors.

If $b \neq b'$ then an acceptor has adopted a higher ballot number from another leader and so the scout terminates with a `Preempted($b'$)` message so that the leader is notified that this ballot has been abandoned.

Commanders

Commanders are spawned and passed a pvalue (b, s, c). A commander spawned with this pvalue operates under the assumption phase 1 has concluded for the ballot b and so the commander will attempt to secure acceptance of this pvalue.

It begins by broadcasting a `phase2a($b,s,c$)` message to the set of acceptors and waits for a response from each. It behaves in a similar pattern to a scout: If the commander already has a majority quorum of responses it discards this response. If not, then the commander examines the (α, b') return type.

If $b = b'$ then the commander adds α to the quorum. If, having added α the commander has secured a majority quorum then the proposal (s, c) has successfully been committed. Hence the commander broadcasts to all replicas a `decision($s,c$)`; otherwise the commander continues to wait for further responses.

If $b \neq b'$ then, just as we had with scouts, the commander will terminate with a `Preempted($b'$)` so the leader can abandon the ballot.

Leaders

Having described the operation of scouts and commanders, it is now necessary to discuss when leaders spawn sub-processes and how they process their responses. Leaders maintain the following important state:

- `ballot_num`: initially $(0, \lambda)$ for leader with identifier λ. This is the ballot over which the leader currently attempts to secure adoption and acceptance of a corresponding pvalue.

- `active`: a boolean value that describes whether the leader is currently in *active* or *passive* mode.

- `proposals`: the list of proposals the leader has received from replicas.

A passive leader has spawned a scout and is waiting for phase 1 of the synod protocol to conclude. The leader initialises in passive mode when it attempts to perform phase 1 with the ballot $(0, \lambda)$. Conceptually we can treat leaders as having receiving three different kinds of messages: proposals from replicas, adoption messages from scouts and preemption messages from scouts or commanders (although of course the implementation of how these are processed differs).

Leaders receive `Propose(`s, c`)` messages from replicas and will add them to their list of proposals, assuming it hasn't already been proposed to a given slot. If the leader is active then it will spawn a commander to attempt to commit this proposal, since a quorum of acceptors will have adopted the leader's current ballot. If the leader is not active then it simply waits until an adoption message is received, in which case it will then spawn commanders for all stored proposals
.
Preemption messages `Preempted(`b`)` occur when a scout receives from an acceptor a ballot number b higher than that they are seeking adoption or acceptance. In this case the commander moves into passive mode and increments the round number of its ballot and spawns a new scout to attempt to secure adoption of this new ballot number.

An `Adopted(`b`,pvals)` message indicates that phase 1 of the protocol has concluded. The leader will in this situation be in passive mode and so enter active mode now that it can begin phase 2 of the protocol. From `pvals` the leader for each slot number the command with the highest corresponding ballot and adds it to their list of `proposals`, removing from `proposals` any that have the same slot number. For each of these proposals the leader spawns a commander sub-process to execute phase 2 of the protocol and seek acceptance for each pvalue.

3.6 Summary

This chapter described the implementation details of the project; beginning with a high level overview of the architecture and then delving into the message system and each role the processes can take and their ability to communicate. Chapter 4 goes on to describe the evaluation of the software that took place.

Chapter 4

Evaluation

With an implementation of Multi–Paxos having been developed it is necessary to evaluate it, both in terms of its ability to perform under the assumptions of the environment in which it operates and also to characterise its performance in a number of typical cases.

4.1 Experimental setup

This section describes the framework upon which simulations were setup to allow for evaluation to take place.

4.1.1 Experimental measurements

Performance evaluation requires the measurement and subsequent analysis of experimental data. There are two key metrics required for evaluating the performance of networked systems – *latency* and *throughput.*

Latency is the time taken elapsed between a client sending a request to the system and receiving a response. It is measured by a client broadcasting a command with a `Nop` operation to the set of replicas and starting a timer. The timer is stopped when the response with the corresponding command identifier is received and the elapsed time is calculated and output to a trace log file.

Throughput is the rate at which the system services requests and hence measured in requests per unit of time. Clients submit requests at a fixed rate (e.g. 10 requests per second) and the number of responses returned per second is measured by clients. This gives a throughput for a given arrival rate and allows for a *peak throughput* measurement to be made. That is, a certain arrival rate of requests will maximise the output throughput which represents the maximal rate at which the system services requests.

4.1.2 Mininet

Mininet is the network simulator that was used for the evaluation of the implementation, running a virtual machine with a Linux guest operating system with 2GB of memory. Mininet ships with a Python API that allows for scripts to run network simulations. Mininet allows for arbitrary Linux processes to be run on any of the virtual hosts. This involved installing an OCaml runtime to execute bytecode compiled on the development machine. The simulation runs Python scripts that performs simulations that proceeds as follows

1. Setup up network topology, populating it with hosts and switches and their interconnects (and the performance parameters)

2. For a given number of each role participating, a JSON configuration file is generated that lists each host IP address and port number

3. Multi–Paxos instances are started via the OCaml runtime on each virtual host with their parameters set (e.g. whether they are clients)

4. The simulation starts and proceeds to execute for a given amount of time

5. The simulation is terminated. Configuration files are deleted and the virtual network is torn down and each process is terminated.

Log files are saved by each process on a shared filesystem with the host machine upon which they are analysed by another Python script which produces plots using matplotlib[1].

Unless otherwise stated each experiment used 1Gbps, 20ms delay links with a 0% packet loss probability[2] and switches with queue sizes of 10^7 packets. These characteristics were chosen to help prevent properties of the network from dominating measurements of system performance. The topology of the networks used is discussed in 4.1.4.

4.1.3 Duelling leaders and the impossibility result

When performing initial experiments on the system it became apparent that systems that have more than one leader would regularly livelock when processing requests from clients. Inspection of log files led to me finding out this was a

---

[1]https://matplotlib.org/

[2]As Cap'n Proto uses TCP the developed system does not suffer failures from packet loss, rather it will simply introduce further delays caused by retransmissions

Figure 4.1: Graph of backoffs

well–known case of *duelling leaders*. This occurs when leaders will attempt to gain adoption of their ballot but are preempted by another leader. The leader will then attempt to gain adoption of another ballot before the other leader's ballot is accepted, causing the other leader to attempt adoption of a higher ballot. The leaders will continually increment their ballots and seek adoption without any given ballot being accept, livelocking the system and leaving it unable to make progress.

There is a well known result called the FLP Impossibility result, outlined in Fischer et al's paper [2] that proves that consensus is never guaranteed in an asynchronous environment that permits crash failures. The problem of duelling leaders is an instance of this problem and so there is no guarantee that consensus can be reached in an asychronous environment.

However, there are methods to vastly reduce the chance that leaders will duel. One simple method is to have a *random backoff* so that each leader will pause their operation for a random amount of time in some range. This increases the chance that a leader will have its ballot accepted before being preempted by another leader. A slightly more sophisticated method is to employ *exponential backoff*. In this case each leader will employ a random backoff within some range and upon being preempted this range is doubled. The backoff time is reduced to an initial value when an adopted message is received by a leader. This reduces the amount of idle waiting that may occur when using a random backoff and no duelling is occurring.

Figure 4.1 shows a number of traces with different backoff strategies. With no backoff, the system halts after 18 requests are processed and makes no further progress. Three traces of random backoffs with different maximum wait times were collected. With greater maximum wait times there are a number of greater sporadic delays in the system as these leaders often have to wait idly even when duelling is not occuring. Making the random timeout too small, as in the 0.25s case, reduces the number of delays but increases the chance leaders will duel for some time before one waits long enough for the other to make progress, resulting in very large delays. The exponential backoff (starting with a value of 0.25s) performs best, with very little delay in the general case and ocassional delays as backoff windows increase in size when leaders do duel.
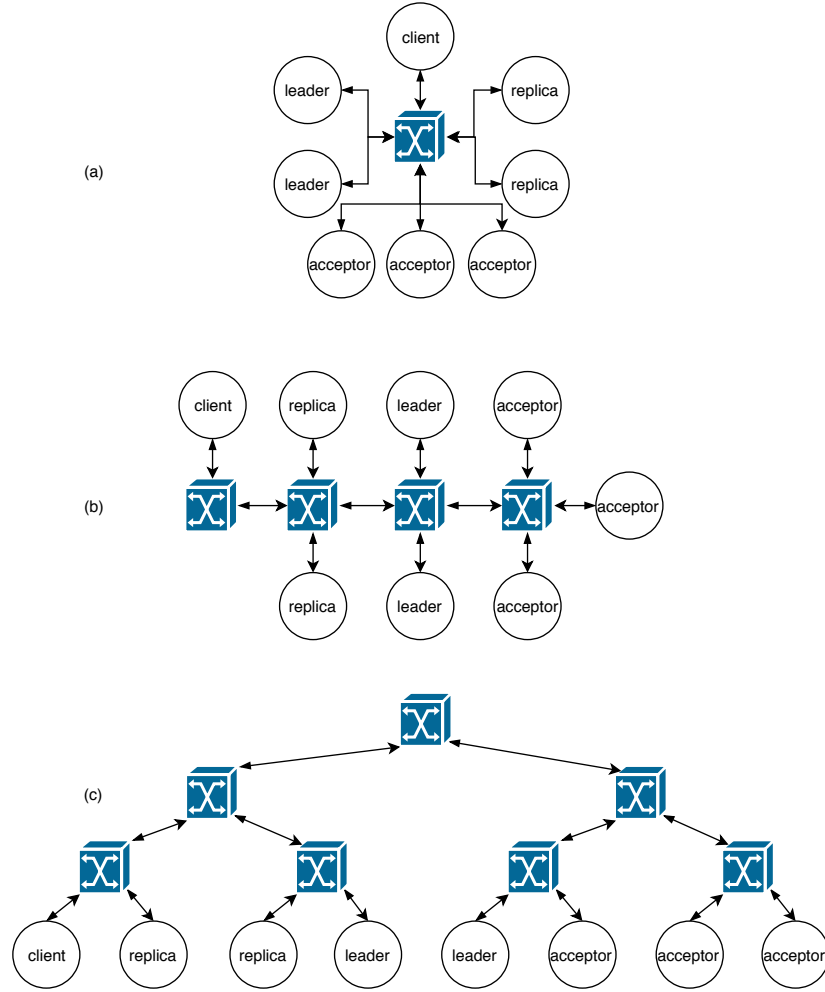
Figure 4.2: Diagrams of different network topologies investigated. Topology (a) represents a star topology, (b) represents a linear topology and (c) a tree topology.

4.1.4   Topologies

An important consideration when evaluating the system on a simulated network is to ensure that the results collected do not reflect the characteristics of the network itself rather than reflecting the characteristics of the system itself. For example we do not wish for delays in queues at switches to dominate the measurements of latency for the system itself. An initial experiment was performed to observe the steady state behaviour of the system with three different network topologies, shown in Figure 4.2. A client submit requests to the replicas every 1.0s and the latency of these requests is measured. Topology (a) represents a star topology wherein each host is connected to a single switch. Topology (b) represents a linear network wherein each process with the same role

Figure 4.3: Graph of latencies over time with different network topologies.

is connected to a single switch and these are each connected together. Topology (c) represents a tree topology with switches arranged in a binary tree structure.

Figure 4.3 shows the traces produced when measuring the steady state latency of each topology. No unexpected behaviour occurs in any of the three schemes. The main difference to note is that the average latency of each trace varies between each topology. This is because the number of hops from each sender and receiver participating in the system varies. For example, every pair of hosts that need to communicate is two hops away in the star topology, 3 hops away in the linear topology and a variable number of hops away in the tree topology. There is no unexpected queuing behaviour resulting in very large delays for any of the given topologies.

4.2 Performance evaluation

The performance evaluation explores two key areas. The first is how the latency and throughput of the system change as the number of processes participating in the protocol grows. The second is how these performance metrics change in the

face of failures.

4.2.1 System size

It is important to characterise the behaviour of the system as the number of processes is varied. For these experiments the number of each role is fixed in a system with 1 client and $f = 1$; that is 2 replicas, 2 leaders, 3 acceptors. Then for each experiment the number of each process was increased and the latency and throughput observed.

Method

For each system size we wish to characterise the average latency and throughput in the steady state. For this experiment the system was arranged in a star topology, with the network parameters described in . When collecting latency measurements, clients submit requests every second for 60 seconds. The first 5 seconds of measurements are discarded since they exhibit non–steady–state behaviour. The experiment produces a collection of latencies $t_1, t_2, t_3, \ldots t_N$. The average latency for this trace is the sample mean of these, that is

$$\bar{t} = \frac{1}{N} \sum_{n=1}^{N} t_n$$

In order to construct a confidence interval, the simulation was repeated $M = 5$ times, producing a collection of average latencies $\bar{t}_1, \bar{t}_2, \bar{t}_3, \ldots \bar{t}_M$. The sample mean of these is computed to give T, the latency averaged over each run of the simulation

$$\bar{T} = \frac{1}{M} \sum_{m=1}^{M} \bar{t}_m$$

and sample variance

$$\sigma^2 = \frac{1}{M-1} \sum_{m=1}^{M} \left(\bar{t}_m - \bar{T} \right)^2$$

This allows us to construct a 95% confidence interval for each experiment, given by

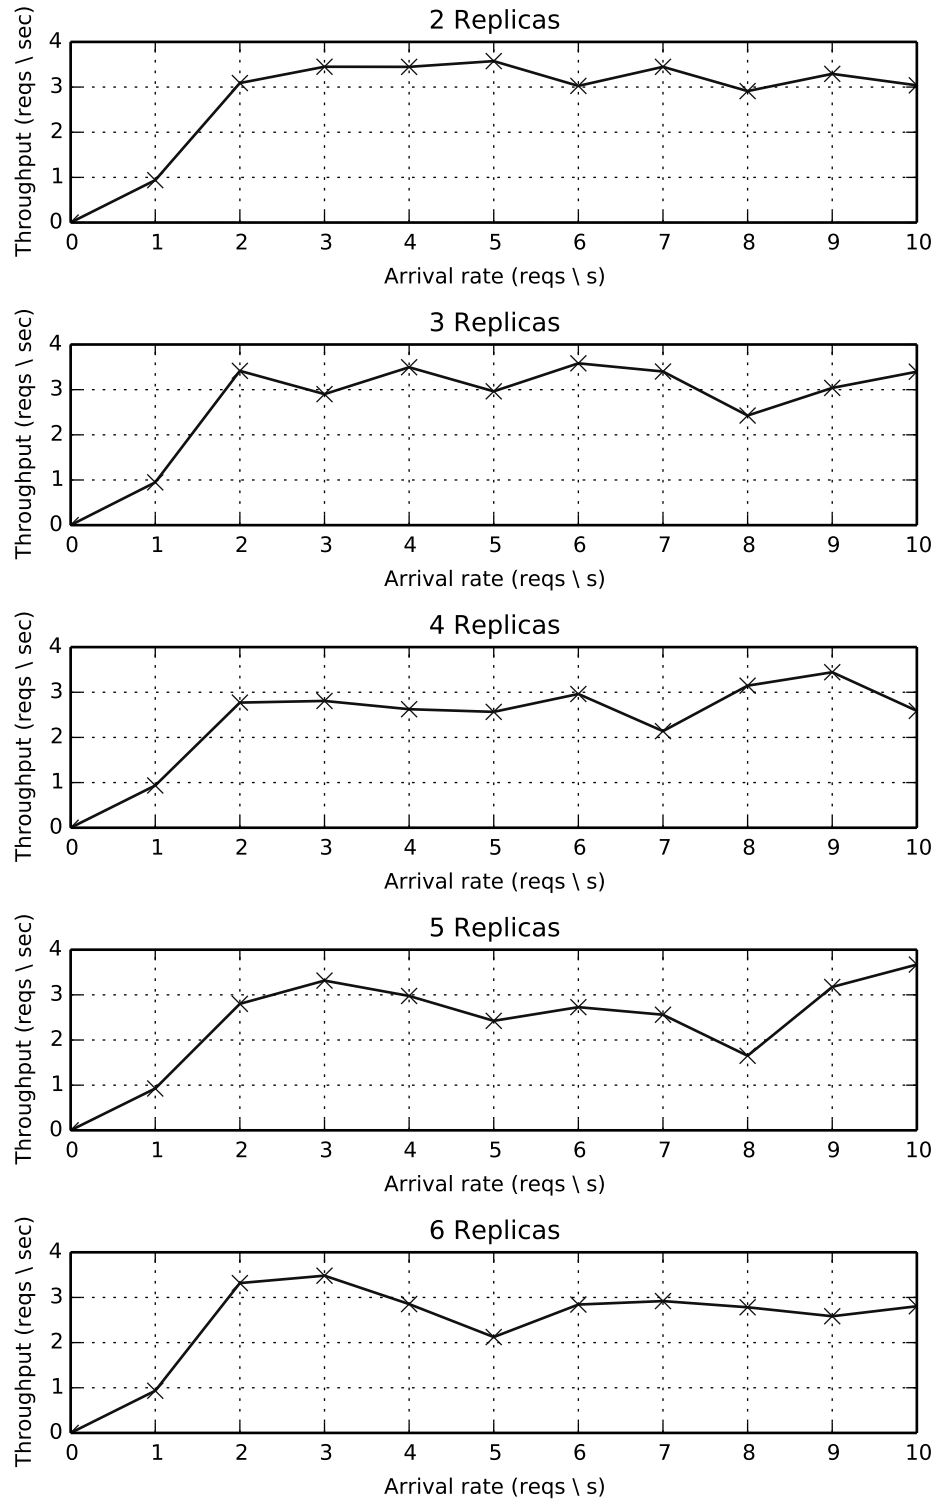$$\left[\bar{T} - \frac{1.96\,\sigma}{\sqrt{M}}, \; \bar{T} + \frac{1.96\,\sigma}{\sqrt{M}} \right]$$

Figure 4.4: Graphs displaying how throughput varies with arrival rate of requests for a number systems with different replica sizes.

The *peak average throughput* was determined for every system size. For each system size, the throughput of the system was measured as the rate of requests from the client, called the arrival rate, was varied. For each arrival rate, the average throughput was computed by dividing the number of responses by the time interval over which they were received.

The results of such an experiment produce a set of graphs like those in Figure 4.4, which contains how these throughputs vary as the number of replicas is increased. Increasing the sending rate of the client causes the throughput to increase; this is because when the system is under–utilised the faster requests are sent the faster they will be processed. This occurs up to a certain sending rate, wherein the system becomes fully utilised and increasing the sending rate does not increase the throughput but instead causes it to oscillate near the limit or drop slightly as the system becomes overloaded.

Results

TODO: Do more runs of the throughput simulations to construct confidence intervals.

TODO: Finish analysing acceptor throughputs.

The results of the experiment are displayed in Figure 4.5. In each case there is no drastic change in latency or throughput of the system as the number of each role is varied.

In the general case most replicas will propose the same serialisation of commands to the leaders which are then de–duplicated. Hence increasing the number of replicas does not introduce additional overhead except to increase the number of messages in the network. This does not incur a performance penalty when the network is not congested. The number of leaders exhibits decreasing latency as the system is increased to 4 leaders and then increases again. This unusual behaviour may be attribute to random variations in the simulations rather than the behaviour of the system; as the confidence intervals of each experiment overlap. This is also reflected in the throughput as the system exhibits unusually high throughput for the case of 4 leaders.
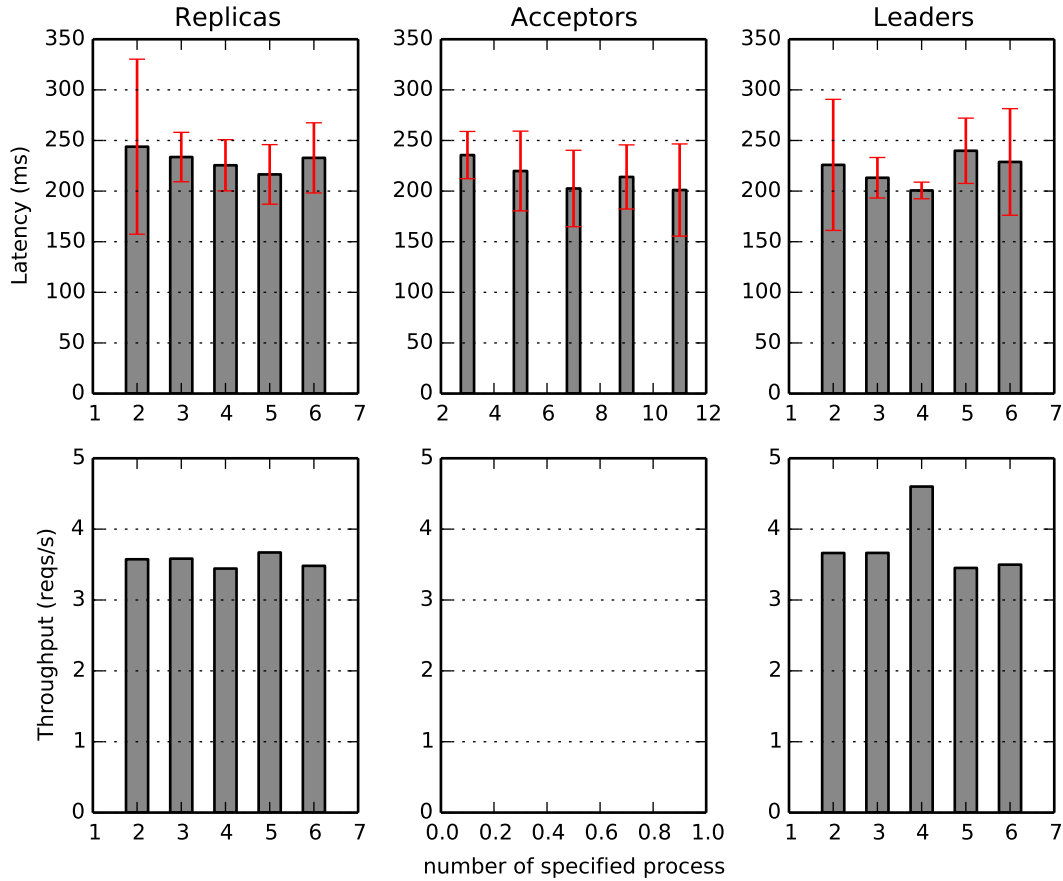
Figure 4.5: Results of latency and throughput as system size was varied for each role separately. For each graph the roles that did not vary were fixed in a system with $f = 1$. Error bars show 95% confidence intervals.

4.2.2 Failure traces

This experiment observes the behaviour of the system when failure of a process is triggered. Each simulation ran for 120 seconds with a client submitting commands at a rate of 1 per second and a system size with $f = 1$. After 60 seconds of simulated time has elapsed a process failure is triggered; this involves the killing of a specific process on a virtual host. The simulation continues to run for another 60 seconds so the behaviour after the failure can be observed. Simulations were performed separately for killing a replica, leader and acceptor.

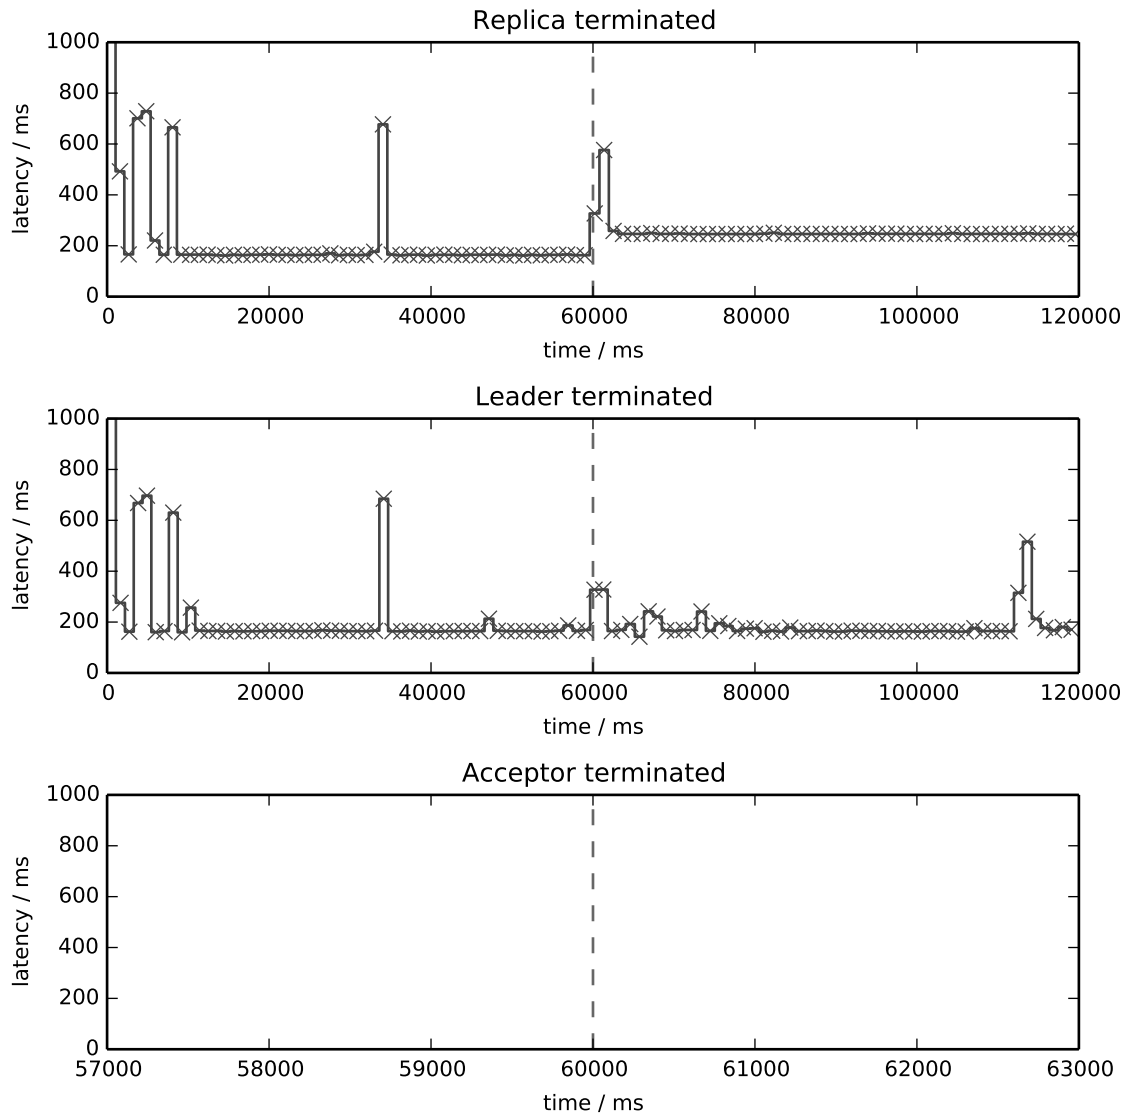The traces produced are shown in Figure 4.6. In each case the system begins as

Figure 4.6: System size

usual with very high latencies before settling into steady state behaviour. In the case of a replica process failing, there is a transient increase in latency before settling down to a latency higher than before the crash. The leader process failing causes an increase in latency, smaller than that of the replica, and causes small transient delays briefly thereafter. Unlike the replica failure the latency returns to that of before the crash.

TODO: Have a look over the acceptor code because it doesn't seem to be behaving as it should. Repeat experiment after.

Chapter 5

Conclusion

The aim of this project was to produce an implementation, in OCaml, of the Multi–Paxos consensus algorithm. This chapter concludes with a consideration of the successes of the project, any limitations that have been discovered, a discussion of possible future work that could be undertaken and closes with any final remarks.

5.1 Successes

The implementation of this project was a success in that an implementation of Multi–Paxos was developed as described in the Requirements section in Chapter 2. A strongly consistent key value store was replicated across a number of replica processes with the consistenty maintained by having their commands serialised by leaders and acceptors. The system tolerated failures as described by the assumptions laid out in 2.1.1.

5.2 Limitations

A limitation encountered over the course of the project was the use of Mininet as a network simulator. Problems arose with linking to external libraries meant the program developed could only be interpreted as OCaml bytecode rather than compiled native code. This led to delays in evaluation and ... . Furthermore, the size of simulations that could be performed was constrained by the resources available to the virtual machine running Mininet. In the future it would be beneficial to run larger simulations with a less constrained network simulator.

Another limitation encountered in the evaluation was in making a performance comparison with LibPaxos. Inexperience with the C language and lack of documentation made it difficult to construct a comparable implementation of the application on top of LibPaxos. Because of this it was not possible to make the desired performance comparsion.

5.3 Future work

The large amount of literature and ongoing research into consensus algorithms presents lots of opportunities for future work on the project. One such opportunity for future work is to modify the project to that of Flexible Paxos and measure the associated performance gains against the original.

Other opportunities could arise from *open sourcing* the project and providing extensions to make it attractive to the wider OCaml community, for example providing compatibility with Mirage OS [8] or creating a GUI interface for visualising log files. A number of optimisations suggested in the literature could be implemented, including introducing *read–only* commands that do not need to be decided on by the synod protocol and reducing the amount of state retained by replicas over time.

5.4 Final remarks

This project has been both insightful and rewarding, providing me with experience designing and implementating a distributed system and exploring the challenges faced when doing so.

Bibliography

[1] Mike Burrows. The chubby lock service for loosely-coupled distributed systems. In *Proceedings of the 7th Symposium on Operating Systems Design and Implementation*, OSDI '06, pages 335–350, Berkeley, CA, USA, 2006. USENIX Association.

[2] Michael J. Fischer, Nancy A. Lynch, and Michael S. Paterson. Impossibility of distributed consensus with one faulty process. *J. ACM*, 32(2):374–382, April 1985.

[3] Heidi Howard, Dahlia Malkhi, and Alexander Spiegelman. Flexible paxos: Quorum intersection revisited. *CoRR*, abs/1608.06696, 2016.

[4] Michael Isard. Autopilot: Automatic data center management. Technical report, April 2007.

[5] Leslie Lamport. The part-time parliament. *ACM Trans. Comput. Syst.*, 16(2):133–169, May 1998.

[6] Leslie Lamport. Paxos made simple. pages 51–58, December 2001.

[7] Leslie Lamport. Fast paxos. *Distributed Computing*, 19:79–103, October 2006.

[8] Anil Madhavapeddy, Richard Mortier, Charalampos Rotsos, David Scott, Balraj Singh, Thomas Gazagnaire, Steven Smith, Steven Hand, and Jon Crowcroft. Unikernels: Library operating systems for the cloud. *SIGPLAN Not.*, 48(4):461–472, March 2013.

[9] Brian M. Oki and Barbara H. Liskov. Viewstamped replication: A new primary copy method to support highly-available distributed systems. In *Proceedings of the Seventh Annual ACM Symposium on Principles of Distributed Computing*, PODC '88, pages 8–17, New York, NY, USA, 1988. ACM.

[10] Diego Ongaro and John Ousterhout. In search of an understandable consensus algorithm. In *Proceedings of the 2014 USENIX Conference on USENIX Annual Technical Conference*, USENIX ATC'14, pages 305–320, Berkeley, CA, USA, 2014. USENIX Association.

[11] Luís Rodrigues and Michel Raynal. Atomic broadcast in asynchronous crash-recovery distributed systems and its use in quorum-based replication. *IEEE Trans. on Knowl. and Data Eng.*, 15(5):1206–1217, September 2003.

[12] Fred B. Schneider. Implementing fault-tolerant services using the state machine approach: A tutorial. *ACM Comput. Surv.*, 22(4):299–319, December 1990.

[13] Robbert Van Renesse and Deniz Altinbuken. Paxos made moderately complex. *ACM Comput. Surv.*, 47(3):42:1–42:36, February 2015.

Appendix A

Cap'n Proto schema file

```
@0x9207445e65eea38d;

interface Message {
  # Interface for RPC messaging system

  struct Command {
    clientId @0 :Data;
    # Id of the client that issued command

    commandId @1 :UInt16;
    # Id of the command being issued

    operation @2 :Text;
   # Encodes the operation that will be applied to state of application
    # The type of this is temporary for now

    clientUri @3 :Text;
  }
  # Structure represents the command sent in RPC

  clientRequest @0 (command :Command) -> ();
  # Method clientRequest is a message sent from the client to a replica
  # The client issues a command and response is returned in another message

  decision @1 (slot_number :UInt16, command :Command) -> ();
  # Replicas receive decision messages sent by a leader
  # Consists of a command and a slot number
  # Slot number is the place slot in which the command has been decided
  # to occupy by the synod protocol

  sendProposal @2 (slot_number :UInt16, command :Command) -> ();
```

```
  # Method sendProposal is a message sent from a replica to a leader.
  # Proposals consists of a command and a slot for which that command
  # is proposed.

  clientResponse @3 (commandId :UInt16, result :Text) -> ();
  # Method clientResponse is a message sent from replica to a client
  # Returns the id of the command and result of issuing it

  phase1 @4 (ballotNumber :Text) -> (result :Text);
  # Method phase1 is a message sent from leader to an acceptor
  # Used to secure adoption of ballot given in argument

  phase2 @5 (pvalue :Text) -> (result :Text);
  # Method phase2 is a message sent from leader to an acceptor
  # Used to secure acceptance of pvalue given in argument
}
```

Appendix B

Projct Proposal

Computer Science Tripos - Part II - Project Proposal

Achieving Distributed Consensus with Paxos

Christopher Jones, Trinity Hall

Originator: Christopher Jones

20th October 2017

Introduction

Paxos is a widely used algorithm that allows consensus to be reached in the context of failure-prone distributed systems, by having a number of processes agree upon a proposed value. A variant of this algorithm, Multi-Paxos, allows for a sequence of values to be agreed upon by electing a leader at the start.

This project wil consist of the implementation of Multi-Paxos in OCaml. To demonstrate an application of distributed consensus, a strongly consistent replicated key-value store will be implemented using this Multi-Paxos implementation. A message passing system, leveraging a RPC library, will be developed that allows nodes running the application to communicate. On top of this messaging functionality, the Multi-Paxos algorithm itself will be implemented.

The performance of this implementation will be evaluated with respect to Lib-Paxos[1], an existing Paxos library that will be used to compare performance to the implementation developed for this project. Testing and evaluation will take place on an emulated network providing a stable and adjustable test environment.

Work to be done

The project breaks down into the following main sections:-

---

[1]http://libpaxos.sourceforge.net

1. Development of a messaging system using an RPC library that allows networked processes to communicate. This will provide the underlying system for message-passing that will be used to implement Multi-Paxos.

2. The development of a test harness that simulates failures, such as processes crashing, restarting and stalling. It will be necessary to ensure these actions can be triggered at specific points in the execution of Multi-Paxos. A network emulator will be used to test possible failure modes such as dropped packets or broken links.

3. The main body of work for the project will consist of the implementation of the core aglorithm. Being distributed in nature, the algorithm consists of sending messages, using the system developed prior, between networked processes. The ability to select random quorums of processes will also need to be included.

4. Evaluation of the algorithm will take place on an emulated network, with varying topologies and a number of failure modes. Tests will be performed to ensure consensus is reached under given assumptions about the network and processes. Performance will be measured in terms of latency and throughput as the number of participating nodes is varied and compared against LibPaxos running under similar test conditions.

Starting point

A large amount of literature is available on Paxos that will be used as a specification of the algorithm that will be developed in this project.

I'm starting the project with no prior knowledge of OCaml and its environment/tools, only some knowledge of Standard ML. OCaml libraries such as Core[2], Async[3] and Cap'n Proto[4] (for RPCs) may be used.

LibPaxos, an existing open source implementation, will be used as a benchmark against which to compare performance in terms of latency and throughput.

---

[2]https://github.com/janestreet/core
[3]https://github.com/janestreet/async
[4]https://capnproto.org

Success criteria

A clear success criterion for this project is that the application that runs Multi-Paxos, the replicated key-value store, is in fact strongly consistent across all replicas.

Paxos operates under a set of assumptions about the network, processes on the network and their respective failure modes. These are assumptions such as packet loss, packet re-ordering and processes that can crash, stall and restart. A success criterion of this project is that given the set of assumptions the implementation of Multi-Paxos achieves consensus and makes progress.

Paxos should be able to make progress if F processes in a network of $2F + 1$ processes fail. This is a key criterion laid out in descriptions of Paxos and as such will be used as a judgement for success - numerous tests to check progress will be conducted, given the simulated failure of up to F processes at a number of points of execution.

The performance of the implementation and the existing LibPaxos library will be compared in terms of latency and throughput as the number of processes on the network is varied. All tests will be performed on an emulated network and will provide a means by which to judge the performance of this implementation against one already used in applications. If the perfomance of the implementation in terms of these metrics is within 30% (an achievable but still desirable performance when compared to a popular library) of that of LibPaxos it will be deemed successful in terms of performance.

Possible extensions

A desirable extension to Mulit-Paxos is Flexible Paxos[3]; a variant of the algorithm that weakens the requirement that quorums in each stage need intersect. This could be evaluated against LibFPaxos[5], a prototypal extension of LibPaxos.

Another possible extension is to implement an interactive application on top of the replicated key-value store, such as a concurrent editor.

---

[5]https://github.com/fpaxos/fpaxos-lib

Work Plan

1. **Michaelmas weeks 3-4** Gain familiarity with OCaml. Prepare build automation, package management, continuous integration and version control. Research Core, Async and Cap'n Proto. *Deadline: 01/11/2017*

2. **Michaelmas weeks 5** Gain familiarity with Mininet, write scripts to produce different network topologies and collect example data. Run a test networked OCaml application on Mininet. Thoroughly research Multi-Paxos algorithm and research possible evaluation strategies. *Deadline: 8/11/2017*

3. **Michaelmas weeks 6** Begin implementation of the project. Develop the key-value store. Integrate the RPC library to allow for processes to communicate. Pass unit tests that confirm this messaging system behaves as expected on Mininet. *Deadline: 15/11/2017*

4. **Michaelmas week 7-8** Define each of the roles nodes play in the algorithm. Begin the implementation of the core algorithm, starting with the leadership election phase. Begin implementing ability to select random quorums, generate unique proposal numbers, prepare and promise requests. *Deadline: 29/11/2017*

5. **Michaelmas vacation** Continue with implementation of aglorithm, finishing phase one. Next complete phase two - implement accept requests / responses. Pass tests to ensure expected functionality. Prepare network environments for evaluation and collect preliminary data. *Deadline: 10/01/2018*

6. **Lent weeks 0-2** Write progress report. Prepare presentation. Continue with evaluation of project by running LibPaxos on Mininet under the same conditions. Collect data on LibPaxos that will be compared to this implementation. *Deadline: 31/01/2018*

7. **Lent weeks 3-4** Finish up any remaining experiments required for evaluation. Calculate confidence intervals of data. Prepare plots for presentation in dissertation. *Deadline: 14/02/2018*

8. **Lent weeks 5-6** If time permits, begin an extension of FPaxos and start testing that under the same conditions. Otherwise, continue any evaluation still outstanding. *Deadline: 28/02/2018*

9. **Lent weeks 7-8** Finish up implementing and evaluating possible extension(s) to the project. Start writing dissertation. *Deadline: 14/03/2018*

10. **Easter vacation** Continue writing dissertation. Complete a draft before the end of the vacation. *Deadline: 18/04/2018*

11. **Easter weeks 0-2** Complete final changes to dissertation. *Deadline: 09/05/2018*

Resource declaration

I will use my own machine (2014 Macbook Air, 1.4 GHz Intel Core i5, 4GB RAM, 128GB SSD) for software development, connected to the University network in order to access online resources. I accept full responsibility for this machine and I have made contingency plans to protect myself against hardware and/or software failure. Should my machine fail I will use the MCS facilities.

Git will be used with Github for version control and regular backups. The development directory will reside in a Google Drive for further backup.

Mininet[6], an open source network emulator, will be used for testing and evaluation. In order to run Mininet on my system, I will use VirtualBox[7].

TODO: Put the proposal's bibliography at the end of the document

---

[6]http://mininet.org
[7]https://www.virtualbox.org/